





# Lecture 0x02:

## Basic Windows Internals

# Agenda for Today: Windows Basics

- Core OS Concepts
- Windows Internals



# The Windows Operating System

# What is an Operating System

A fancy resource manager.

It abstracts away the managing of physical devices (CPU, Memory, I/O devices...etc)

# Operating System Kernel

Responsible for servicing requests and translating them into instructions compatible with computer hardware including the CPU, Memory, I/O devices...etc

Often provides an interface for userland applications to interact indirectly with resources

# The Windows Operating System

- Proprietary Operating System Developed by Microsoft
- Multi Arch: 32/64 bit, preemptive multitasking operating system
- Primary flavors are Client, Server
  - Others include IOT, CE, Xbox, Some defunct targeting mobile...etc

We will focus on Windows (10) Client/Server



# Devices that run Windows

- Desktop computers
- Servers
- Laptops
- Tablets
- XBox (technically a fork of Windows 2000)

Different devices have different considerations (battery, screen size, memory, storage ...etc)

# (Some) Windows Design Principles

- Extensibility
- Portability
- Reliability
- Performance
- Compatibility
- Security
- \$\$\$\$

# Extensibility

- The bulk of system services are provided by the Executive
- The Windows OS is technically a monolithic OS
  - I.e., One driver can mess with another driver. Kernel memory is shared
- However, it separates userland execution environments into “subsystems”
- Each subsystem has the ability to execute a type of application
- This provides a modular setup where changes to execution environments don’t necessarily necessitate changes to the executive

# Portability

- Windows will run on a wide variety of systems.
- The bulk of the OS is written in C and C++
  - The rest is bundles of code used for processor specific instructions (Eg Arm vs intel/amd)
- Platform dependent code is implemented in the Hardware Abstraction Layer (HAL)

This is less magical than it was 10 years ago as most modern OSes support variety in hardware/architecture

# Reliability


- User applications should not be able to corrupt the operating system

Windows audits most drivers and has rigorous standards for which drivers are officially signed

That said, malicious drivers can muck around with kernel memory. There are countermeasures for this. (eg Driver Guard)

# Performance

- Should be fast.



# Compatibility

Backwards compatibility and extending of existing technology

This is partially why the Windows API is so stable

# Security

- Kernel/User separation
- User separation
- Process (and associated resources) isolation



# Windows Kernel

- Windows is an Object based kernel
- Resources are called objects, and can only be directly accessed by the kernel
- To indirectly interact with an object, a *handle* is used

# Kernel Objects

- Kernel object (KOs): a single run-time instance of a statically defined object type
- Object types are system-defined data types.
- Each object type has its own attributes and functions to interact with it
- For example, an object of type *process* is an instance of a process object.
- A *file* object is an instance of a file. Note, *file!= a thing on disk*

# Objects and Handles

A handle is an abstract reference to an object. This could be an actual pointer to the object, or a reference to a per-process GUID that references an object

This allows us to abstract away direct management of objects in memory, and instead work with with references. **This is a security control. If something goes wrong in kernel space, you get a BSOD.**

APIs are used to interact with system resources, share resources among processes, and protect resources from unauthorized access.

# Windows System Architecture

- User Processes
- Subsystem DLLs
- NTDLL.dll
- Service Processes
- Executive
- Kernel
- Device Drivers
- Win32k.sys
- (sometimes) Hyper-V

# Interacting with the Windows OS

- Windows API functions (win32): Documented, callable functions in the Windows API.
  - For example, MessageBox, CreateFile, GetMessage
- Native system services (sys calls): Undocumented (officially) underlying services in the OS that are callable from user mode. For example
  - NtAllocateVirtualMemory is the internal service used for VirtualAlloc
  - NtCreateUserProcess is the internal service used by CreateProcess
- Direct Syscalls
- Other: (WinRT, COM,...etc)
- Kernel support functions: functions inside the Windows OS that can only be called from in kernel mode

# Windows Executables and Shared Libraries

# Executable File Formats

- Bundles of machine code and associated data needed to run a program
- Usually requires an OS to *load* the executable
- Code is organized according to a convention that the programmer and the kernel agree on.

It is just a convention though!

# PE File Format Basic Definitions and Concepts

- Portable Executable (PE) is an executable file format used by Windows NT
- It contains information about code to execute, and how it should be executed
- In this discussion, we will use an open source tool PE-Bear to look at the structure of a PE file.



# PE File Format

- PE file format is used for both userland and kernel mode executables
  - Userland: file.exe, file.dll, file.obj
  - Kernel mode: driver.sys, ntoskrnl.exe
- PE is based on the Common Object File Format (COFF).
- PE format is not architecture specific (hence “portable”)
  - Note this means the format can be used across multiple different architectures. The target architecture is still specified inside of the PE though
- Data is grouped together in blocks called *sections*, identified by *headers*

# PE: Libraries

- Windows shared libraries are called Dynamically Linked Libraries
- They are PE files with a special characteristic set
- They can export functions from their code
- Other PEs can load these libraries and access the exported code
- This allows for modular programs

We will talk more about this next lecture when we review linking.

# PE

PEs are composed of sections and headers

Sections are data/code

Headers contain information about how to load the PE, and where the data is, and how to process it

# Tools

In this lecture, we will use PE-Bear to explore the PE file format.

As a sample, lets use Calc.exe (64 bit)

Run `$path = Get-Command calc.exe` to find the path to calc.exe on your machine

Run `PE-Bear.exe $path.Source` (in powershell)

# Calc.exe

File Settings View Compare Info

calc.exe

- DOS Header
- DOS stub
- NT Headers
  - Signature
  - File Header
  - Optional Headers
- Section Headers
  - .text → EP = C70
  - .rdata
  - .data
  - .pdata
  - .rsrc
  - .reloc

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
C70	48	83	EC	28	E8	2B	FA	FF	FF	48	83	C4	28	E9	7E	FD
C80	FF	FF	CC	CC	CC	CC	CC	CC	CC	CC	CC	CC	CC	CC	CC	CC
C90	48	83	EC	28	E8	0B	01	81	38	63	73	6D	E0	75	23	83
CA0	78	18	04	75	1D	8B	48	20	8D	81	E0	FA	6C	E6	83	F8
CB0	02	76	08	81	F9	00	40	99	01	75	07	FF	15	5F	09	00
CC0	00	CC	33	C0	48	83	C4	28	C3	CC	CC	CC	CC	CC	CC	CC
CD0	48	83	EC	28	E8	0D	B5	FF	FF	FF	FF	15	9F	08	00	
CE0	00	33	C0	48	83	C4	28	C3	CC	CC	CC	CC	CC	CC	FF	25
CF0	4C	09	00	00	CC	CC	CC	CC	CC	CC	CC	CC	48	83	EC	18
D00	33	D2	48	8D	41	FF	48	83	F8	FD	77	3C	B8	4D	5A	00

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
H	.	i	(	+	+	+	+	+	+	+	+	+	+	+	+	+
y	y	i	i	i	i	i	i	i	i	i	i	i	i	i	i	i
H	.	i	(	H	.	.	.	.	.	.	.	.	.	.	.	.
x	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.
H	.	i	(	H	.	.	.	.	.	.	.	.	.	.	.	.
.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.
L	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.	.
3	0	H	.	.	.	.	.	.	.	.	.	.	.	.	.	.

Disasm: .text General DOS Hdr Rich Hdr File Hdr Optional Hdr Section Hdrs Imports Resources Exception BaseReloc Debug LoadConfig Hint

	Hex	Disasm	Hint
1870	4883EC28	SUB RSP, 0x28	
1874	E82BFAFFFF	CALL 0x1400012A4	
1879	4883C428	ADD RSP, 0x28	
187D	E97EFDFFFF	JMP 0x140001600	
1882	CC	INT3	
1883	CC	INT3	
1884	CC	INT3	
1885	CC	INT3	
1886	CC	INT3	
1887	CC	INT3	
1888	CC	INT3	
1889	CC	INT3	
188A	CC	INT3	
188B	CC	INT3	
188C	CC	INT3	
188D	CC	INT3	
188E	CC	INT3	
188F	CC	INT3	
1890	4883EC28	SUB RSP, 0x28	
1894	488B01	MOV RAX, QWORD PTR [RCX]	
1897	813863736DE0	CMP DWORD PTR [RAX], 0xE06D7363	
189D	7523	JNE SHORT 0x1400018C2	
189F	83781804	CMP DWORD PTR [RAX + 0x18], 4	
18A3	751D	JNE SHORT 0x1400018C2	
18A5	8B4820	MOV ECX, DWORD PTR [RAX + 0x20]	
18A8	8D81E0FA6CE6	LEA EAX, [RCX - 0x19930520]	
18AE	83F802	CMP EAX, 2	
18B1	7608	JBE SHORT 0x1400018BB	
18B3	81F909409901	CMP ECX, 0x1994000	
18B9	7507	JNE SHORT 0x1400018C2	
18BB	FF155F090000	CALL QWORD PTR [RIP + 0x95F]	[msvcrt.dll].?terminate@@YAXXZ
18C1	CC	INT3	

# DOS Header

- The DOS header contains the magic bytes MZ that identify it as a PE
- The final entry (offset 0x3c referenced as `->e_lfanew`) is the offset the of NT Headers



# NT Headers

- Signatures
- File Header
- Optional Header

The screenshot displays the Windows File Explorer interface for the file `calc.exe`. The left pane shows the file's structure, with the **NT Headers** section expanded and highlighted by a red rectangle. The right pane shows the hex dump of the file's content, with the address `E8` highlighted by a red rectangle.

Address	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
E8	50	45	00	00	64	86	06	00	10	C4	40	03	00	00	00	00
F8	00	00	00	00	F0	00	22	00	0B	02	0E	14	00	0C	00	00
108	00	62	00	00	00	00	00	00	70	18	00	00	00	10	00	00
118	00	00	00	40	01	00	00	00	00	10	00	00	00	02	00	00
128	0A	00	00	00	0A	00	00	00	0A	00	00	00	00	00	00	00

# Signature

- Usually 4 bytes containing
- “PE\0\0”
- For our purposes, it is only used to verify the file format.

A screenshot of a hex editor window. The interface includes a toolbar at the top with icons for navigation and editing. The main area is a table with three columns: offset (hex), hex data, and ASCII data. The offset column has labels E8, F8, 108, and 110. The hex data column shows values at these offsets. At offset E8, the first four bytes (50 45 00 00) are highlighted in yellow. The ASCII data column shows the string 'PE..' at offset E8, with the 'P' and 'E' characters highlighted in yellow. The background of the hex editor is dark grey.

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
E8	50	45	00	00	64	86	06	00	10	C4	40	00	00	00	00	00
F8	00	00	00	00	F0	00	22	00	0B	02	0E	10	00	00	00	00
108	00	62	00	00	00	00	00	00	70	18	00	00	00	00	00	00
110	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00



# File Headers

Following the Signature, we have the File Headers. This gives us

- The number of sections (NumberOfSections)
- Whether or not we have a DLL/EXE (Characteristics)
- The Compilation timestamp
- A pointer to a symbol table if one exists

Disasm	General	DOS Hdr	Rich Hdr	File Hdr	Optional Hdr	Section Hdrs	Imports
Offset	Name	Value	Meaning				
EC	Machine	8664	AMD64 (K8)				
EE	Sections Count	6	6				
F0	Time Date Stamp	340c410	Friday, 24.09.1971 16:02:24 UTC				
F4	Ptr to Symbol Table	0	0				
F8	Num. of Symbols	0	0				
FC	Size of OptionalHeader f0	240					
▼ FE	Characteristics	22					

# Optional Headers

I don't know why it is listed as optional. I don't think a PE can run without this section (but I could be wrong?)

The optional headers contains most of the data required to load PE

Specifically, values found here are used to build the *Import Address Table*, and perform *Base Relocations*

Offset	Name	Value	Value
100	Magic	20B	NT64
102	Linker Ver. (Major)	E	
103	Linker Ver. (Minor)	14	
104	Size of Code	C00	
108	Size of Initialized Data	6200	
10C	Size of Uninitialized Data	0	
110	Entry Point	1870	
114	Base of Code	1000	
118	Image Base	140000000	
120	Section Alignment	1000	
124	File Alignment	200	
128	OS Ver. (Major)	A	
12A	OS Ver. (Minor)	0	
12C	Image Ver. (Major)	A	
12E	Image Ver. (Minor)	0	
130	Subsystem Ver. (Major)	A	
132	Subsystem Ver. Minor)	0	
134	Win32 Version Value	0	
138	Size of Image	8000	
13C	Size of Headers	400	
140	Checksum	14163	
144	Subsystem	2	Windows GUI
▼ 146	DLL Characteristics	C160	
		40	DLL can move
		100	Image is NX compatible
		4000	Guard
		8000	TerminalServer aware
148	Size of Stack Reserve	80000	
150	Size of Stack Commit	2000	
158	Size of Heap Reserve	100000	
160	Size of Heap Commit	1000	

# Optional Headers (pt 1)

- Magic: Architecture of image
- Entry Point: Relative virtual address (RVA) from the Base Address
- Image Base: (preferred) Base address: Where in memory the PE “prefers” to be loaded. If the location is unavailable, the Image needs to be relocated

Offset	Name	Value	Value
100	Magic	20B	NT64
102	Linker Ver. (Major)	E	
103	Linker Ver. (Minor)	14	
104	Size of Code	C00	
108	Size of Initialized Data	6200	
10C	Size of Uninitialized Data	0	
110	Entry Point	1870	
114	Base of Code	1000	
118	Image Base	140000000	
120	Section Alignment	1000	
124	File Alignment	200	
128	OS Ver. (Major)	A	
12A	OS Ver. (Minor)	0	
12C	Image Ver. (Major)	A	
12E	Image Ver. (Minor)	0	
130	Subsystem Ver. (Major)	A	
132	Subsystem Ver. Minor	0	
134	Win32 Version Value	0	
138	Size of Image	B000	
13C	Size of Headers	400	
140	Checksum	14163	
144	Subsystem	2	Windows GUI
146	DLL Characteristics	C160	
		40	DLL can move
		100	Image is NX compatible
		4000	Guard
		8000	TerminalServer aware
148	Size of Stack Reserve	80000	
150	Size of Stack Commit	2000	
158	Size of Heap Reserve	100000	
160	Size of Heap Commit	1000	

## Optional Headers (pt 2)

- **SizeOfImage**: the virtual size of the image
- **SizeOfHeaders**: the size of the headers
- **DLLCharacteristics**: flags including knowledge of hardening features such as ASLR/ CFG...etc. Not super important for us other than assuming knowledge of ASLR.

Offset	Name	Value	Value
100	Magic	20B	NT64
102	Linker Ver. (Major)	E	
103	Linker Ver. (Minor)	14	
104	Size of Code	C00	
108	Size of Initialized Data	6200	
10C	Size of Uninitialized Data	0	
110	Entry Point	1870	
114	Base of Code	1000	
118	Image Base	140000000	
120	Section Alignment	1000	
124	File Alignment	200	
128	OS Ver. (Major)	A	
12A	OS Ver. (Minor)	0	
12C	Image Ver. (Major)	A	
12E	Image Ver. (Minor)	0	
130	Subsystem Ver. (Major)	A	
132	Subsystem Ver. Minor)	0	
134	Win32 Version Value	0	
138	Size of Image	8000	
13C	Size of Headers	400	
140	Checksum	14163	
144	Subsystem	2	Windows GUI
146	DLL Characteristics	C160	
		40	DLL can move
		100	Image is NX compatible
		4000	Guard
		8000	TerminalServer aware
148	Size of Stack Reserve	80000	
150	Size of Stack Commit	2000	
158	Size of Heap Reserve	100000	
160	Size of Heap Commit	1000	

Size of the Image in bytes, as well as the headers

Subsystem (are we a console application or GUI?)

Windows GUI

DLL can move  
Image is NX compatible  
Guard  
TerminalServer aware

# Sections

.text Executable code (machine instructions)

.data: global variables

.rdata: read only global variables

# Dynamically Linked Libraries:

- Refresh: What is a DLL?

# DLL

A PE with with DLL characteristic field set.

Usually it has exported functions which can be referenced by code outside of the DLL

The Windows API is implemented in a handful of DLLs that export specific functions

# Win32 API

We will mostly leverage documented functions from the Windows API

The function definitions are well documented

Reading that documentation however, is a skill that must be learned

Sometimes, we need more control over what we are trying to accomplish, and will leverage undocumented functions stored in NTDLL.dll



# NtDLL.dll

Implements the Windows Native API. This is the lowest layer of code that is still Userland code.

It is used to communicate with the kernel for system call invocation.

NtDLL also implements the Heap Manager, the (executable) Image loader and some of userland thread pools. Every process loads this DLL in the same location in memory!

# Kernel32.dll

Contains (more or less) the same functionality as NtDLL!

It exposes basic operations such as memory management, input/output (I/O) operations, process and thread creation, and synchronization functions

It can be thought of as a compatibility layer, as it almost always calls directly into NTDLL.dll

This is to maintain backwards compatibility- where the Win32 API rarely changes, but the Native API changes from release to release.

# 32bit vs 64bit

This class will focus on 64bit executables (Intel x86 64)

When developing code that needs to run on either a 32bit or 64bit systems, you need to take care when assuming the size of various types. For example, type sizes vary across 32bit and 64bit architectures and you need to take care when defining them.

# Processes

# Processes

Nothing in userland is executed outside of the context of a process.

You don't "run a processes"

You run threads

Processes are containers, and there is no such thing (to my knowledge) as code running outside of a process

# Threads

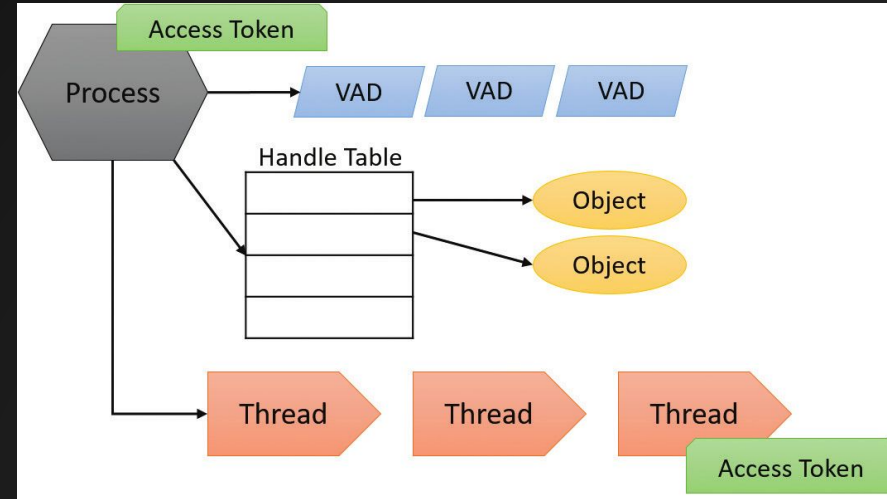
Unit of execution contained within a process

- Actual entity that executes code

More on threads in later. For now, we will only deal with processes that have a single executing thread.

# Process Resources

- Uniquely identified by a Process ID
- Contains a security context in the form of a Process Access Token
- Maintains a table of handles to objects
- Has a *private* virtual address space
- $\geq 1$  Thread (possibly with its own token)



# Virtual Memory

An abstraction layer around physical memory

Each process gets its own private virtual memory

$2^{32}$  bytes worth on 32 bit (possibly more with some extensions)

$2^{64}$  bytes worth on 64 bit


The kernel lies about how much space there really is

Think of a virtual address space as a giant, contiguous array of bytes



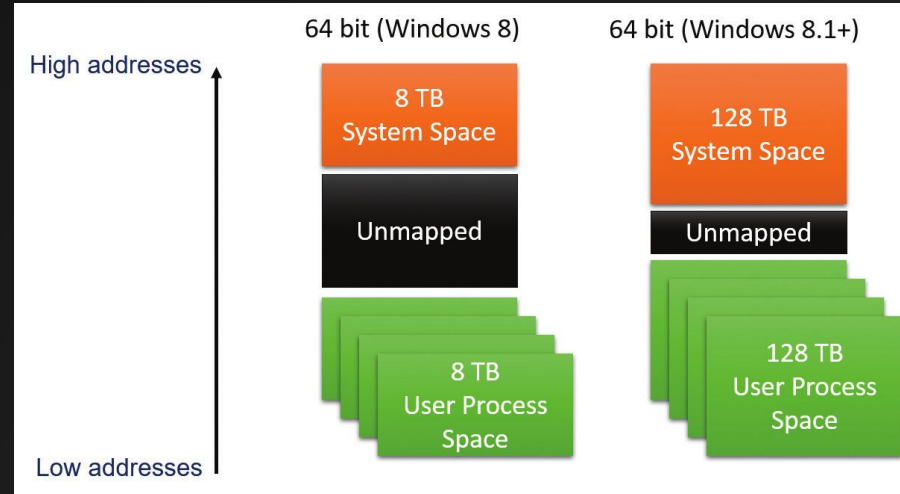
# Process Virtual Memory

- Windows implements a linear Virtual Address space
- Provides a logical interface for a process to interact with indirectly with physical memory



# Memory Layout

- Each processes gets its own *Virtual Address Space*
- The Windows OS divides Virtual Address space into two portions: kernel and Userland\*
- Memory here grows “downward” from higher address spaces to lower address spaces
- On x64, the upper half is reserved for kernel space and the bottom half is reserved for processes



# Memory

Virtual: A Block of raw memory

Stack: technically can be managed by the programmer but is usually managed by the processes and the compiler

Heap Memory: managed by the programmer with a special data structure called a Heap (Basically a priority queue)

# Stack: what is it good for?

- Local variables
- Passing arguments to a function
- Returning values from a function
- Function invocation/ABI

# Stack Memory

- Data structure built on top of our Virtual Address Space that allows us to Push, and pop values from a stack
- Whenever we invoke a function with *call* a new *stack frame* is created. We call this the function call stack
- This is a contiguous chunk of memory that acts as a working space for a function's duration

# Stack Memory

- Stack memory is temporary: once the function returns the memory is reclaimed
- Still error prone but is typically safer and faster than Heap allocations
- A programmer usually does not need to worry about managing stack memory
- We need to know how much is needed at compile time!

# Heap Allocation

Allows for dynamic allocation sizes.

Recall we don't need to know how much memory is needed at compile time

Heap memory is managed by a heap allocator.

The programmer allocates with a call to malloc and frees the memory with a call to free

If the programmer forgets to call free, that memory is now unusable until the process terminates and we have introduced a *memory leak*

# Heap Memory

- Memory that can be allocated/deallocated at runtime
- Used for data whose size is not known at compile time
- Slower than stack





# Process Creation

Somewhat complicated. We will simplify it for this class, than dive into it next class after we talk about handles.

- Kernel opens the image (executable file) and verifies it is the correct format
- The kernel creates a new process kernel object and a thread kernel object
- The kernel maps the image to an address space, as well as ntdll.dll
  - Note this gets mapped to just about every type of process
- The creator process notifies Windows subsystem process (Csrss.exe) that a new process and thread have been created
- From the kernel's perspective, the process is created at this point
- Some magic happens, imports are resolved and after all the required LLs are declared, we reach the entry point and the program starts

# Basic Information of a Process

- Name: Usually the executable name. This is NOT a unique identifier
- Process ID (PID: Unique ID of a process. PIDS are reused after a process terminates
- Status: Running, Suspended, Not Responding
- Username: the user who is running the process. It also includes the primary token that holds the security context for the user
- Session ID: Session number under which the process executes. Session 0 is for system processes and services. Session 1 and higher are used for interactive logins.